

---

# Recommender systems with TensorFlow

Martin Jakomin  
Jan Hartman

---

---

# About us

Data scientists



Dr. Martin Jakomin



Jan Hartman



---

# What do we do?

- We develop and maintain the “brains” behind the purchasing of online ad space
  - Choosing the right ad <-> recommender system!
  - We use TensorFlow for click-through rate prediction (230+ million predictions per second)
  - Involved in many other interesting things (conversion, bid win prediction, bid price control... )
-

---

# Agenda, part 1

- Recommender systems paradigms and algorithms
    - Content-based filtering:
      - Item feature extraction
      - Similarity measures and distances
    - Collaborative filtering:
      - Memory based (user-user, item-item, hybrid)
      - Basic algorithms (nearest neighbours, ....)
      - Model based - latent factors
      - Matrix Factorization
    - Deep learning for RS:
      - Deep Neural CF
      - Autoencoders
      - DeepFM
  - Recommender systems evaluation and metrics
  - A short case study - YouTube recommendations
-

---

# Agenda, part 2

- **TensorFlow**: basic concepts, tensors, autodiff
- Keras high-level API: layers, models, fitting & evaluation
- Matrix factorization implementation
- Deep autoencoder implementation
- Bonus challenge

Repo link: [github.com/janhartman/recsystf](https://github.com/janhartman/recsystf)

---

---

# Introduction

- We often make choices in life without sufficient personal experiences, thus we heavily rely on recommendations.
- Recommender systems are tools that assist and augment this natural decision making process.



---

# Introduction

- RS are large-scale machine learning and knowledge discovery tools aimed at providing **personalized** recommendations.
  - Examples: videos, movies and songs, web pages, articles, ads, books, social media posts, products, services ...
  - In terms of machine learning problems: estimation of user ratings, picking topN (ranking), picking new unfamiliar items (novelty, diversity).
-

---

# Paradigms And Algorithms

Typically we split RS into three categories:


- ***Content-based filtering*** (where the user is recommended items similar to the ones he preferred in the past),
- ***Collaborative filtering*** (where the user is recommended items that people with similar tastes and preferences liked in the past),
- ***Hybrid approaches*** (that combine collaborative and content-based methods).



---

# Content-based Filtering

- Recommendations are based on user choices made in the past (e.g. previous purchases), combined with rated items **descriptions**.
- The main idea: if you like a particular item in the past you will also like a **similar** item in the future.



---

# Content-based Filtering

- Similarity between objects is therefore the basis for recommending new items.
  - We can view this as a user-specific classification problem of user's likes and dislikes based on an item's features.
  - CBF is best suited to situations where there is abundant data on items, but not for the users.
-

---


# Content-based Filtering

- |  |   |   |
|--|---|---|
| ★ Easy to implement                      | ↓ | Difficult information extraction                      |
| ★ Transparent, Explainable               | ↓ | Limited content analysis                              |
| ★ User independent (no user data needed) | ↓ | Overspecialization (only recommending similar things) |
| ★ Scalable                               | ↓ | User cold-start                                       |
| ★ No item cold-start                     |   |   |
-

---

# Content-based Filtering

- The recommendation process is usually performed in 3 steps:
  - Content Analysis (extracting features),
  - Profile learning (extracting user preferences),
  - Filtering (extract relevant items).



---

# Item Similarity

- Item representation:
  - Structured,
  - Unstructured (vector space models, TF-IDF),
  - Mixed.

$$w_{i,j} = tf_{i,j} \times \log\left(\frac{N}{df_i}\right)$$

$tf_{i,j}$  = number of occurrences of  $i$  in  $j$

$df_i$  = number of documents containing  $i$

$N$  = total number of documents

---

---

# CBF - Recommending

- Recommend simply by returning similar items to those which the user has rated in the past.
  - Weighting can be applied in case of real-number ratings.
  - More sophisticated approaches can use Bayesian Classifiers, cluster analysis, decision trees, and neural networks in order to estimate the probability that the user is going to like the item.
-

# Similarity Measures

- Euclidean distance:

$$\sqrt{\sum_{i=1}^n (q_i - p_i)^2}.$$

- Manhattan distance:

$$\sum_{i=1}^n |p_i - q_i|,$$

- Minkowski distance


$$D(X, Y) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}}$$

- Jaccard distance

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}.$$

- Cosine distance

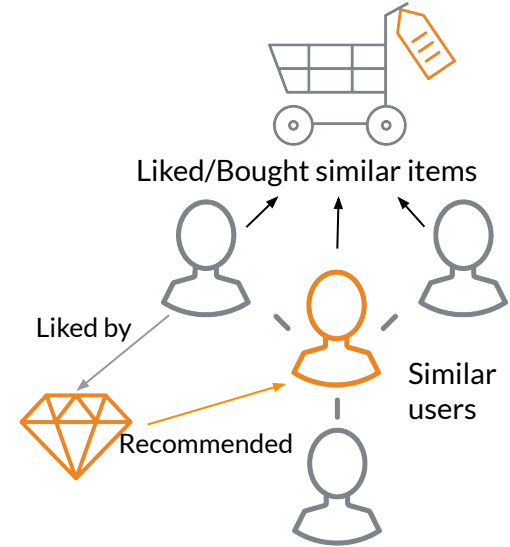
$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$



---

# Collaborative Filtering

- Recommendations are usually based on user's neighbours choices made in the past (of similar items).
- The main idea: people who agreed in the past will **agree** in the future and will like similar kinds of items.





---

# Collaborative Filtering

- Oppose to CBF, we generally only use past rating information (though, other information can be incorporated into models).
  - No “feature extraction”, models can learn abstract (latent) features on fly.
  - This allows for serendipitous recommendations (items can be completely different from what was seen and liked in the past).
-

---

# Collaborative Filtering

- ★ Superior modelling and accuracy
  - ★ Serendipity
  - ★ Efficient and scalable
  - ★ Justifiable and transparent\*
  - ★ Robust and stable
- ↓ Black box models\*
  - ↓ User and Item cold-start
  - ↓ Can be biased

\* Depends on the model

---

---

# Memory-based CF


- Most basic approach where ratings are calculated using simple statistics and aggregations upon stored data.
- Very simplistic, but can often capture local associations in the data.
- Generally split into:
  - ***user-based*** (where users are recommended items from their neighbouring users with similar rating patterns),
  - ***item-based*** (where users are recommended items based on their ratings of the set of similar items).

---

# User-based CF

- Predict a rating for a user  $u$  for an item  $i$  using the ratings for the item  $i$  by users most similar to  $u$ .
- Weighted kNN (regression or classification).

$$\hat{r}_{ui} = \frac{\sum_{v \in \mathcal{N}_i(u)} w_{uv} r_{vi}}{\sum_{v \in \mathcal{N}_i(u)} |w_{uv}|}$$




---

# Item-based CF

- Predict a rating for a user  $u$  for an item  $i$  using the ratings by the user  $u$  for items most similar to  $i$ .
- Weighted kNN (regression or classification).

$$\hat{r}_{ui} = \frac{\sum_{j \in \mathcal{N}_u(i)} w_{ij} r_{uj}}{\sum_{j \in \mathcal{N}_u(i)} |w_{ij}|}$$



---

# Memory-based CF

- Memory-based CF methods are simplistic, transparent (explainable), efficient and allow for serendipitous predictions.
  - We can further extend MCF with various normalization (centering and biases), choosing different similarity metrics (Euclidean, Cosine, ...) or by applying various pre-filtering and post-filtering stages (neighbourhood selection), speedups (locality-sensitive hashing), ...
  - Finally, we can combine user-based and item-based approaches using Hybrid models.
-

---

# Model-based CF

- Contrary, model-based approaches rather use ratings in order to learn a **predictive model** (which is then used to make new predictions).
  - Goal of modelling user-item interactions with factors representing their latent characteristics.
  - These factors are inferred solely from the ratings and might represent obvious dimensions (e.g. movie genres) or something completely uninterpretable to us.
-

---

# Matrix Factorization


- Essential tool in machine learning, broadly used for dimensionality reduction, compression, clustering, classification, ...
  - Also, one of the most popular and most used CF model.
  - Simple and efficient way to model explicit data and various other side information (implicit ratings, biases, temporal effects, etc.).
-



# Matrix Factorization

- User-Item ratings can be simply represented with a matrix.
- Recommendations now becomes a matrix completion problem.
- We can solve this by factorizing this matrix into a product of (two) smaller matrices.

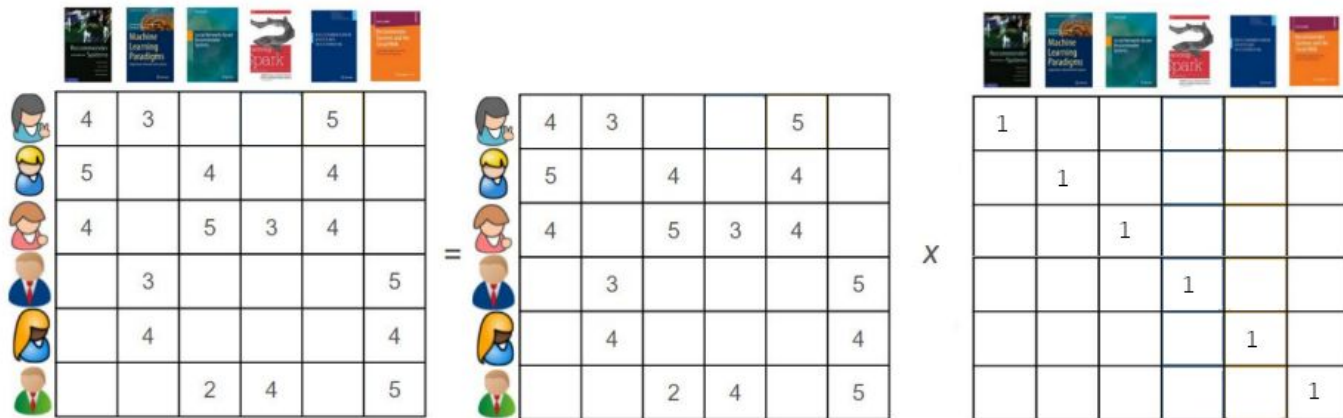
$X$   
 $n \times m$



4	3		?	5	
5		4		4	
4		5	3	4	
	3				5
	4				4
		2	4		5

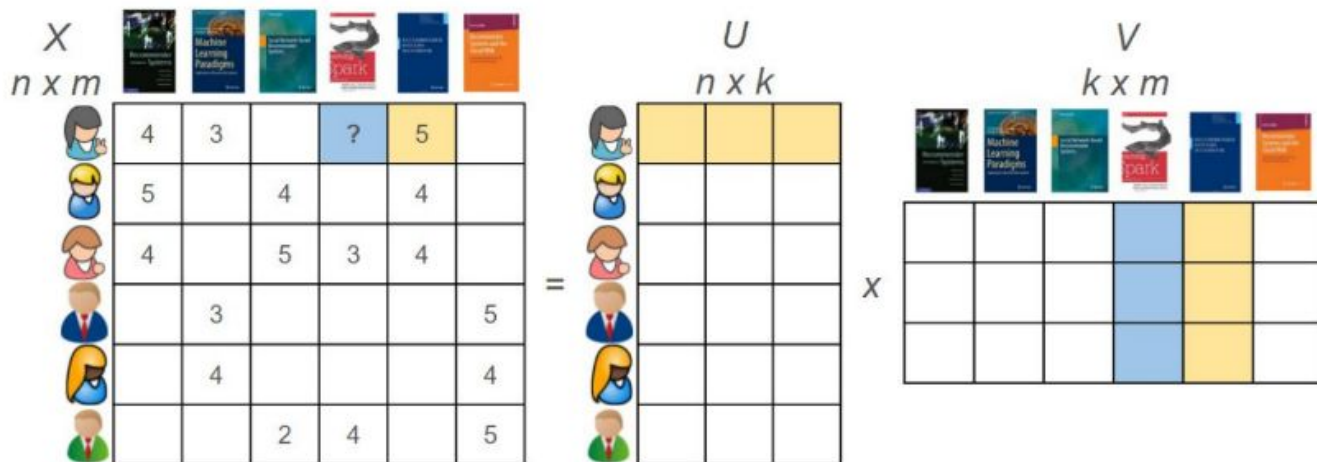
# Matrix Factorization - Intuition

- A trivial factorization:  $X = X * I$

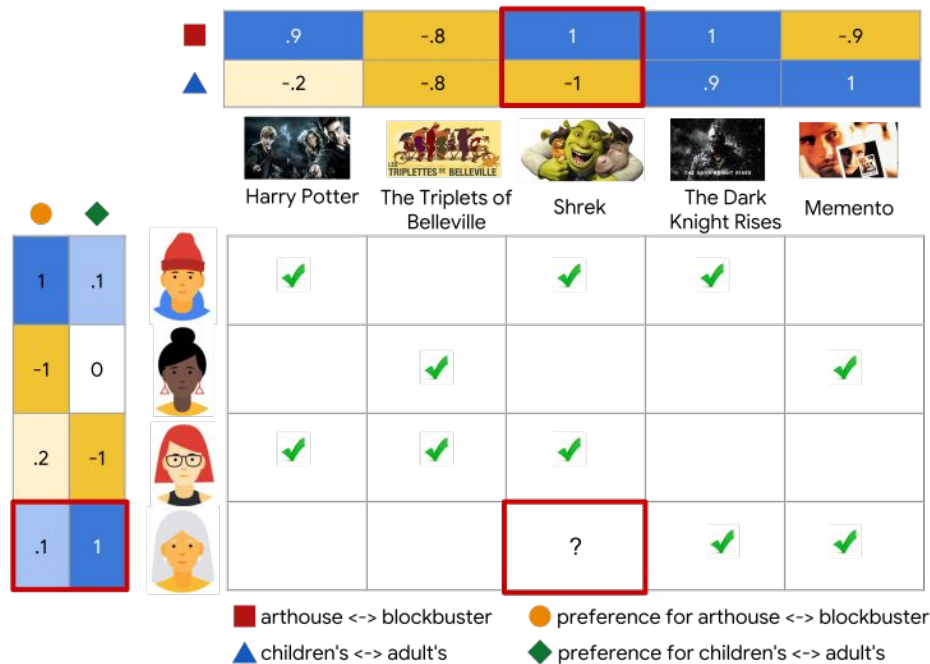
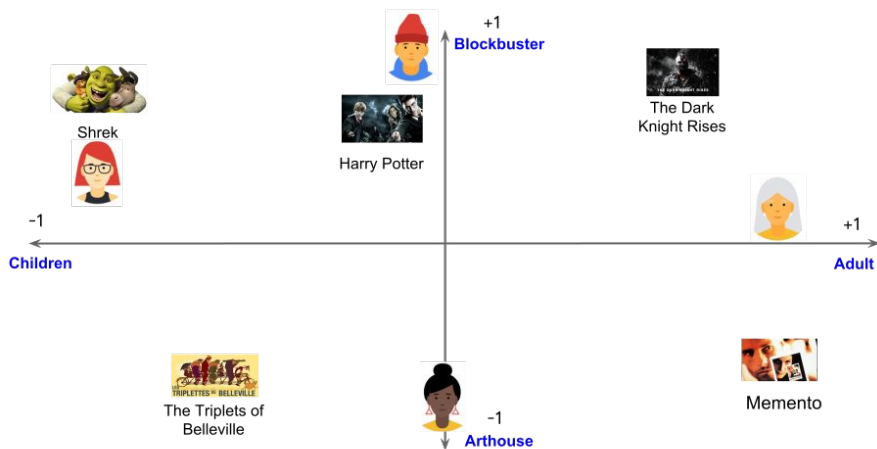


# Matrix Factorization

- Low-Rank factorization:  $X = U * V$



# Matrix Factorization



# Matrix Factorization



---

# Matrix Factorization

- Solve matrix factorization problem:  $X \approx \hat{X} = UV^T$
  - Predict with dot product (linear combinations) between latent user matrix and latent item matrix.
  - There are many possible factorizations - how to find the best one?
  - Define the matrix completion as an optimization problem.
-

---

# Matrix Factorization

- Define the objective (loss) function.
- One of the most common is the squared loss (Frobenius

norm): 
$$||X - \hat{X}||_F^2 = ||X - UV^T||_F^2 = \sum_{i=1}^n \sum_{j=1}^m (x_{ij} - u_i v_j^T)^2$$

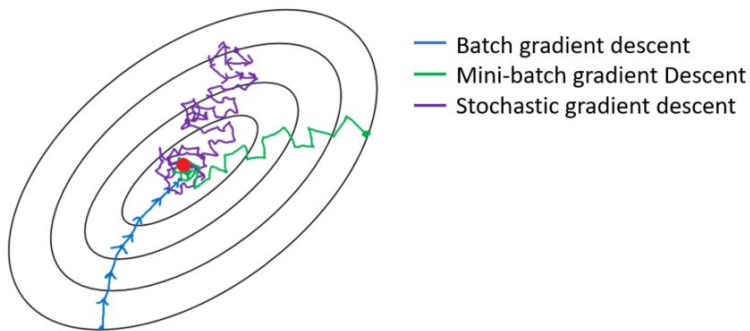
- Problem of huge sparsity (typically over 99%).
- We can use weighted matrix factorization:

$$W \circ X \approx W \circ \hat{X} = W \circ (UV^T)$$

---

# Matrix Factorization

- This kind of optimization problems are usually solved using various gradient descent methods: GD, SGD, MUR, PGD, CD ...
- Alternatively, we can use ALS (convex optimization).
- With SGD we can omit the weighted part and optimize only on the known (non-zero) ratings.





---

# Matrix Factorization - SGD

$$L = \sum_{u,i} (r_{ui} - (\mu + b_u + b_i + \mathbf{x}_u^\top \cdot \mathbf{y}_i))^2 + \lambda_{xb} \sum_u \|b_u\|^2 + \lambda_{yb} \sum_i \|b_i\|^2 + \lambda_{xf} \sum_u \|\mathbf{x}_u\|^2 + \lambda_{yf} \sum_i \|\mathbf{y}_i\|^2$$

$$\frac{\partial L}{\partial b_u} = 2(r_{ui} - (\mu + b_u + b_i + \mathbf{x}_u^\top \cdot \mathbf{y}_i))(-1) + 2\lambda_{xb}b_u$$

$$\frac{\partial L}{\partial b_u} = 2(e_{ui})(-1) + 2\lambda_{xb}b_u$$

$$\frac{\partial L}{\partial b_u} = -e_{ui} + \lambda_{xb}b_u$$

$$b_u \leftarrow b_u + \eta(e_{ui} - \lambda_{xb}b_u)$$

$$b_i \leftarrow b_i + \eta(e_{ui} - \lambda_{yb}b_i)$$

$$\mathbf{x}_u \leftarrow \mathbf{x}_u + \eta(e_{ui}\mathbf{y}_i - \lambda_{xf}\mathbf{x}_u)$$

$$\mathbf{y}_i \leftarrow \mathbf{y}_i + \eta(e_{ui}\mathbf{x}_u - \lambda_{yf}\mathbf{y}_i)$$

---

---

# Matrix Factorization

- MF allows for a simple inclusion of various side information (like implicit ratings, tags), constraints (penalties and rewards), regularization, biases and temporal dynamics.
- Contextual information about the rating process (circumstances) can also be incorporated - CARS:
  - Contextual pre-filtering, post-filtering and modelling
  - $R : User \times Item \times Context \rightarrow Rating$
  - Tensor factorization

---

# Matrix Factorization

- Constrained Matrix Factorization:
  - NMF (only additive linear combinations induce parts-based representation)
  - Sparse MF (constrained number of combinations)
  - Orthogonal MF (uniqueness, clustering)
- Data Fusion

$$\begin{array}{c} R \\ \begin{bmatrix} * & R_{12} & \dots & R_{1r} \\ R_{21} & * & \dots & R_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ R_{r1} & R_{r2} & \dots & * \end{bmatrix} \end{array} \approx \begin{array}{c} G \\ \begin{bmatrix} G_1 & & & \\ & G_2 & & \\ & & \ddots & \\ & & & G_r \end{bmatrix} \end{array} \begin{array}{c} S \\ \begin{bmatrix} * & S_{12} & \dots & S_{1r} \\ S_{21} & * & \dots & S_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ S_{r1} & S_{r2} & \dots & * \end{bmatrix} \end{array} \begin{array}{c} G^T \\ \begin{bmatrix} G_1^T & & & \\ & G_2^T & & \\ & & \ddots & \\ & & & G_r^T \end{bmatrix} \end{array}$$

---

# Deep Learning for RS

- Nowadays, deep learning is everywhere - RS are not an exception.
- Data is becoming more abundant and available.
- Main Idea - the user/item interaction may be non-linear.
- Deep models and Hybrids are SOTA.

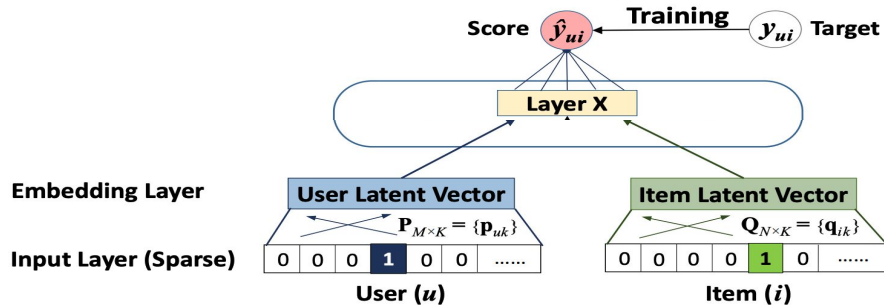


---

# Deep RS

- |   |   |   |
|---|---|---|
| ★ Non-linear transformations                        | ↓ | Black box models - bad interpretability |
| ★ Representation learning and better generalization | ↓ | Data requirement and sparseness         |
| ★ Sequence modelling (RNN)                          | ↓ | Extensive Hyperparameter tuning         |
| ★ Flexibility                                       |   |   |
| ★ Superior accuracy                                 |   |   |
| ★ Unified representation and data fusion            |   |   |
-

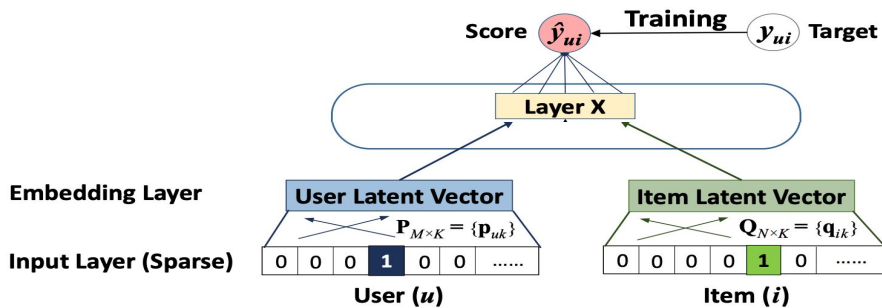
# Deep Neural CF

- User and Item embeddings (with one-hot input)
- Dense layers and single output




# Deep Neural CF

- This is a generalization of the matrix factorization



# Deep Neural CF


- We can add multiple dense layers





# Deep Neural CF


- Or combine both models



# Deep Factorization Machines

- Main idea: Learn both, low-order (FM) and high-order (DNN) feature interactions.


	User				Item				Categories				History				Quantity
$x^1$	1	0	0	...	1	0	0	...	1	0	1	...	1	0	1	...	2
$x^2$	1	0	0	...	0	1	0	...	0	2	1	...	0	0	1	...	4
$x^3$	0	1	0	...	1	0	0	...	3	0	14	...	1	0	0	...	5
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$x^n$	0	0	0	...	0	0	1	...	0	1	5	...	0	0	1	...	1



---


# Autoencoders

- Generalization of the PCA
- Aims to reconstruct the input on the output level



# Deep Autoencoders

- Item-based
- User-based: denoising autoencoder



---


# Evaluating Recommender Systems

- Inherently difficult, same dataset is used for learning and evaluation, high sparsity, bias and outliers
  - How to measure success? User satisfaction is hard to define!
  - How to optimize for novelty and serendipity?
  - In general, we try to optimize the pre-defined loss function (which does not guarantee best predictions or best user satisfactions).
-

---

# Evaluating Recommender Systems

- We commonly use a holdout set (sub-matrix, random sample) on which we apply predictive accuracy metrics, such as RMSE, MAE, ...
- For topN recommendations we typically use coverage measures, such as precision and recall (and F1).



---

# YouTube - A short case study

- Largest platform for creating, sharing and discovering video content
- And also one of the largest scale recommender system in industry
- Main Goal: Maximize the time users spend on the platform (and thus maximize the number of served adds)
- Vast majority of views come from recommendations (more than 70 %)
- However, YouTube is not fully transparent how their algorithms work

---

# History

- In the past YouTube recommended videos according to click-rate, but this incentivized users to create click-baity videos
  - YouTube then changed their strategy to recommend videos based on overall watch time
  - Emphasis on user satisfaction and quality watch time (hard to measure)
  - Search results are also personalized
  - Offline metrics, such as precision, recall, ranking loss are widely used, however the most important metric is the live A/B testing
-



---

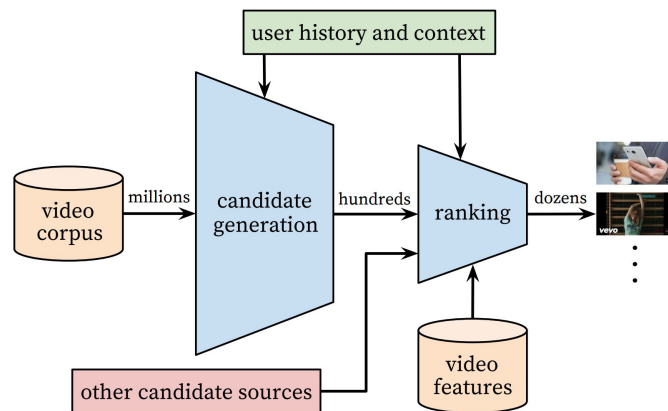
# Deep Neural Networks for Youtube Recommendations

- Used matrix factorization in the past
- One of first to use deep neural network for recommendations
- Built upon TensorFlow
- Two-stage design (with two deep neural networks), one for candidate generation and one for ranking
- This two stage approach allows for recommendation from millions of different videos

---

# The two-stage design

- The candidate generation uses collaborative filtering in order to provide broad personalization (using various features) with a high precision
- Ranking network scores each video according to the desired objective function and thus provides fine-level representation among candidates with a high recall



---

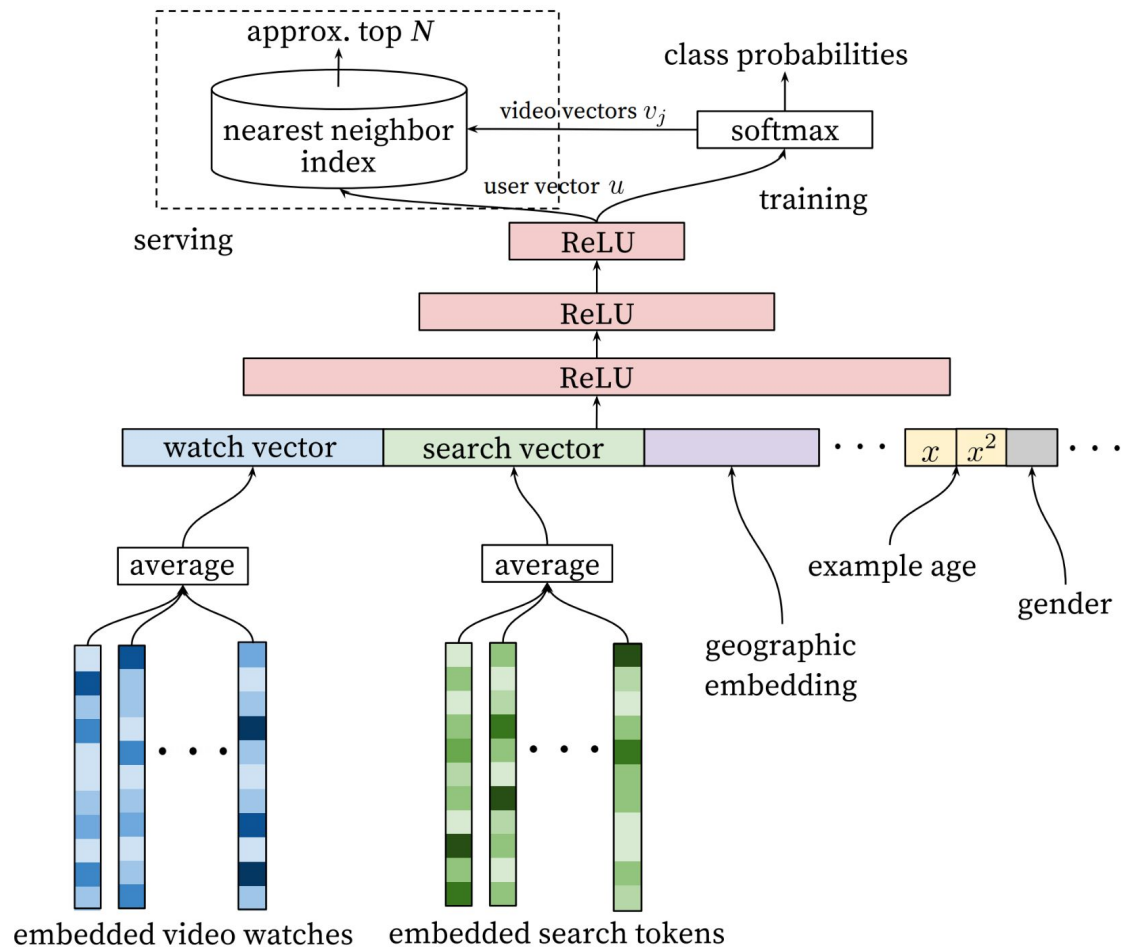
# Candidate generation

- Neural networks as a generalization of matrix factorization (for embedding user's behaviours)
  - Recommendation is viewed as an extreme multiclass (1M) classification
  - Explicit data is very sparse, but implicit (such as video views) is more abundant
  - User watch history is represented by a sequence of sparse video IDs (which is then embedded)
  - Arbitrary continuous and categorical features can be easily added to the model
-

---

# Candidate generation

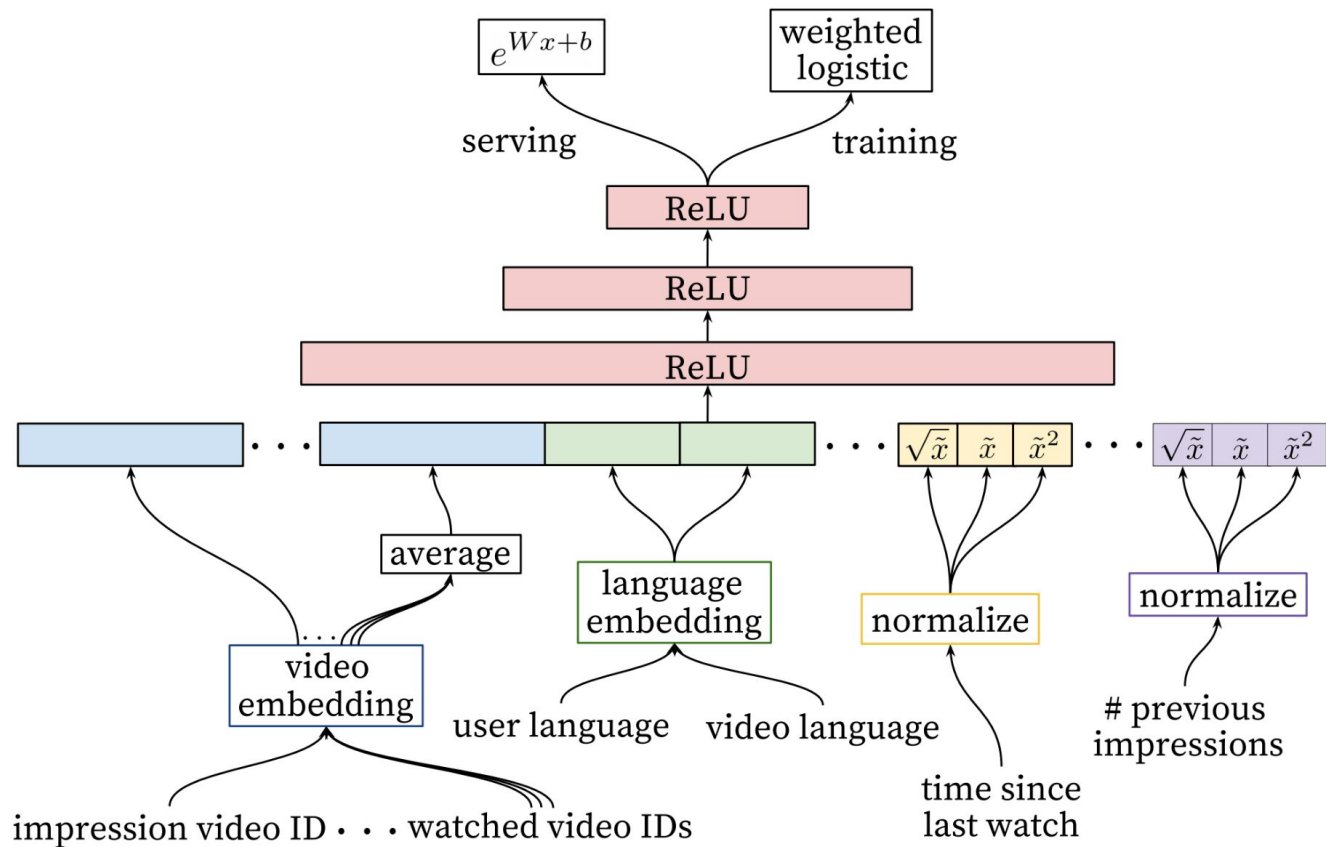
- Age of the video is very important feature in order to recommend “fresh” examples
  - Candidates are found by nearest neighbour search from the last layer of the network
  - Limited number of training example per user (to avoid heavy raters)
  - Adding features and layers significantly improves the accuracy
-



---

# Ranking

- A similar deep neural network is used as for candidate selection
  - More video features (such as thumbnails, etc.) are used
  - Feature engineering is very important (although deep learning alleviates this problem)
  - Shared embedding for similar features and normalization of continuous features
  - Adding features and layers significantly improves the accuracy
-



---

# Add-ons

- Such design is still not optimal (often only relevant/similar items are recommended and users can get bored)
  - Google then introduced REINFORCE algorithm (reinforcement learning) with goal to maximize users' engagement over time by predicting which recommendations would expand their tastes
  - YouTube is constantly updating their machine learning algorithms and has multiple A/B tests in production
-



---

# Conclusion - Key Takeaways

- Offline metrics are a good indicator, but do not always map perfectly to the online setting
  - A/B testing in production is a must
  - Continuous improvements and optimizations
  - Although deep networks help, manual feature engineering still helps
  - Adding features and layers usually improve the accuracy (but has higher infrastructure costs)
  - Optimize for user satisfaction
-

---

# TensorFlow

---



---

# What is TensorFlow?

- *"An end-to-end open source machine learning platform"*
  - At its core: general purpose ML framework, focused on DL
  - Core implemented in C / C++ / Python, wrappers for Java / Go / Swift
  - Peripherals: TF Serving (microservice for serving models), TensorBoard (model visualization / profiling), etc.
-

---

# Pros & cons of TF

- + A lot of implemented operations and models
  - + Optimized for speed, CPU, GPU support
  - + Ecosystem (users, libraries, end-to-end platform) - most widely used DL framework
  - Complexity
  - Poor documentation + awkward API in some places
-

---

# Code time

- Repository: [github.com/janhartman/recsystf](https://github.com/janhartman/recsystf)
  - Option 1: run the notebooks on Google Colab
    - Colab: <https://colab.research.google.com/>
    - Click GitHub and enter *janhartman/recsystf*
    - All dependencies are preinstalled
    - The colab notebooks are stored on your GDrive
  - Option 2: run the notebooks locally
    - Clone the repo and open it
    - Create virtualenv and install dependencies
    - Run *jupyter-lab*
-

---

# Stuff we left out

- Saving and loading models
  - TF Serving, production usage
  - TF in other programming languages
  - Custom op creation
  - Estimators and tf.data pipelines
-

---

# Sources (TF part)

- [TensorFlow guide](#)
  - [Deep Learning With Keras: Recommender Systems](#)
  - [Collaborative Filtering for Movie Recommendations](#)
  - [NVIDIA/DeepRecommender: Deep learning for recommender systems](#)
  - [An implementation of DeepRecommender in Tensorflow & Keras](#)
-

---

# Challenge

---



---

# Data

- Data: [Goodreads book ratings](#), filtered
  - 84k train, 21k test examples
  - *train.csv* & *test.csv*: user ID, book name & rating
  - Bonus: additional book data in *books.csv* (match by book name) - try to find a way to add it to the model
  - *books.csv*: authors, publish date, publisher, rating, number of reviews, number of pages
-

---

# Challenge

- Goal: implement a recommender system in TF to accurately predict a user's rating of a book
  - Use mean squared error as your loss function & a validation dataset
  - Play around with different algorithms, add extra book data into the model...
  - Try to achieve the lowest MSE on the test set (mine: 0.84)
  - Learn TF by experimenting!
-

---

# Conclusion

---



---

# We are hiring!

- Always looking for interesting candidates
  - Various positions are open
  - Hard and interesting problems
  - Big data
  - [zemanta.workable.com](https://zemanta.workable.com)
-



---

# Thank you for your attention!

**Martin Jakomin**

[mjakomin@outbrain.com](mailto:mjakomin@outbrain.com)

[linkedin.com/in/martin-jakomin](https://www.linkedin.com/in/martin-jakomin)

**Jan Hartman**

[jhartman@outbrain.com](mailto:jhartman@outbrain.com)

[linkedin.com/in/jan-hartman](https://www.linkedin.com/in/jan-hartman)

---